


(a) accuracy against training scale



Error bars are SD across training seeds (ddof=1), NOT confidence intervals: neither series has a clustered interval. The two series do not share a seed set, three against five, so this is a system-level juxtaposition and not a matched paired comparison. The accuracy axis does not start at zero.

(b) does the reasoner reduce the multi-step deficit?

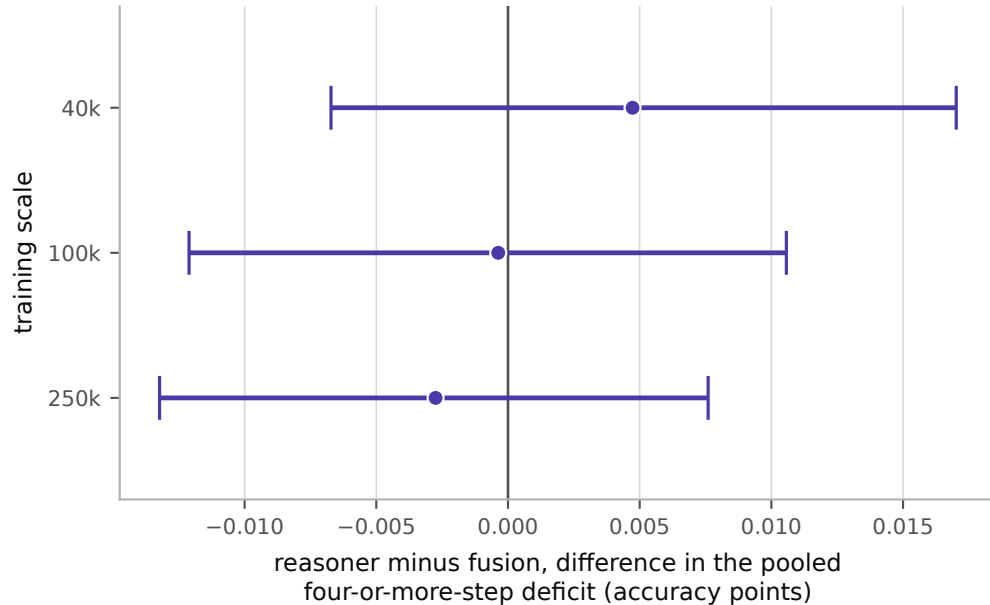


Image-clustered 95% intervals on 7,714 development questions over 768 represented images. All three intervals contain zero: an absence of a detected difference, not evidence of equivalence. Input granularity, architecture and parameter count change together.