

Representation-quality sensitivity at train_40k




Image-clustered 95% intervals on 7,714 development questions over 768 represented images, five training seeds.

Encoder, embedding width (512 against 768) and head input width move together: representation quality is NOT isolated, and this does not show it is the sole or the main bottleneck. No clustered contrast exists at 250k; that direction is reported in VE1-TAB-A3.