




Accuracy against measured end-to-end serial latency, by measurement node. The two panels are never merged.



Every timed row on this node has a VE-0-resolved accuracy pairing, so there are no latency-only rows.



All points are END_TO_END_SERIAL timing, fp32 inference, batch size 1, warm median over three passes. No interval is drawn: a warm median is not an estimate carrying sampling uncertainty. Cached and head-only figures never appear on this axis. Computational efficiency, measured as latency.

The two nodes are never merged into one frontier: the shared fusion control differs by more than the pre-registered tolerance and no adjustment factor is applied. Contextual positioning only, never a leaderboard: the compact VLMs differ in training history, answer support and output format.