


Frozen small-language-model pretraining, question side and answer side

A cross-experiment synthesis of three separately run experiments, not one factorial interface experiment



All rows use the pinned G21 normalised scorer on 7,714 development questions over 768 represented images; intervals are image-clustered and uncorrected for multiplicity. Every row is pretrained minus its architecture-matched random control at the SAME model size. This is a cross-experiment synthesis of three separately run experiments, not one factorial interface experiment: checkpoint selection differs, E8A taking the best development epoch while E8B and E10 use the frozen fixed-22 rule, and the 135M-to-360M step moves trainable capacity too, 21,343,808 against 21,540,800 parameters. An interval containing zero is an absence of a detected effect and establishes NEITHER equivalence NOR the absence of an effect. E10 does not reproduce E8B's directional negative 40k result. E8B B3 at 40k has a seed sd of 0.02795; see VE1-TAB-04.