





Accuracy against measured end-to-end serial latency, by measurement node. The two panels are never merged.



All points are END_TO_END_SERIAL timing, fp32 inference, batch size 1, warm median over three passes. No interval is drawn: a warm median is not an estimate carrying sampling uncertainty. Cached and head-only figures never appear on this axis. Computational efficiency, measured as latency.

The two nodes are never merged into one frontier: the shared fusion control differs by more than the pre-registered tolerance and no adjustment factor is applied. Contextual positioning only, never a leaderboard: the compact VLMs differ in training history, answer support and output format.

Arm codes: A0p = CLIP question encoder; A1 = frozen pretrained SmoLM2-135M question encoder; B1 = lightweight classifier readout; B2 = random SmoLM2-135M readout; B3 = pretrained SmoLM2-135M readout; R2 = trie-constrained greedy generation; R3 = free greedy generation with a fixed token cap; reasoner = latent-query reasoner. 'free' and 'constrained' on a SmoIVLM point are free greedy and answer-support-constrained generation.