





Single node otter159. These points are never merged with the otter155 frontier. CONTEXTUAL POSITIONING ONLY, never a leaderboard or a fair-protocol superiority claim: training history, multimodal pretraining, answer support and output format all differ, and SmolVLM-256M's text backbone is the Instruct checkpoint where E8A and E8B use base. No comparison is made against published official GQA scores in either direction.

Arm codes: A0p = CLIP question encoder; A1 = frozen pretrained SmolLM2-135M question encoder; B1 = lightweight classifier readout; B2 = random SmolLM2-135M readout; B3 = pretrained SmolLM2-135M readout; R2 = trie-constrained greedy generation; R3 = free greedy generation with a fixed token cap; reasoner = latent-query reasoner. 'free' and 'constrained' on a SmolVLM point are free greedy and answer-support-constrained generation.